Data Mining Coursework

Junaid Ali Rasheed

March 26, 2018

Abstract

In this report, we provide a solution for a Portuguese banking institution looking to determine whether a client will make a term deposit or not. The institution provided us with two datasets, a training dataset containing 36,169 instances, and a testing dataset containing 9042 instances. A variety of data exploration and visualisation techniques were used to examine the dataset which determined that we could not predict whether or not a client will subscribe to a term deposit using just two attributes. Quite a few of the numerical attributes provided were negatively skewed. Attribute selection techniques were used to reduce the original sixteen attributes into two reduced datasets; one with three attributes and another with eight attributes. We then ran a number of machine learning algorithms on these datasets and performed various experiments in order to optimise the parameters used for these algorithms. The J4.8 model with the minimum number of objects parameter set to 10 performed best and was estimated to be 90.6% accurate (using 10-fold cross-validation). These models were also assessed under an unequal cost scenario, where misclassifying a client who will subscribe to a term deposit costs 10 times more than misclassifying a client who will. The same J4.8 model also performed best in the unequal cost scenario, with a cost of 10561. The model was then ran on the test dataset which contained 9042 clients. The default rule on the test dataset had an accuracy of 87.8% and a cost 7941. The model achieved an accuracy of 90.3% in the equal cost scenario and a cost of 2641 in the unequal cost scenario.

1 Introduction

We have been provided with a dataset which contains information about customers who were targets of direct marketing campaigns of a Portuguese banking institution. Our task is to develop models on this dataset to determine whether a customer will make a term deposit or not. We will use equal and unequal costs to develop the model.

2 Data Exploration

The dataset provided by the banking institution is split into two files; **cworkTrain.arff** and **cworkPredict.arff**. **cworkTrain.arff** will be used to train the models. These models will then be evaluated on **cworkPredict.arff**

2.1 The Training Dataset

The training dataset is **cworkTrain.arff**. Opening the file in a text editor gives us a brief description of each attribute in the dataset. The attributes in the dataset are described in Table 1.

Name	ID	Description
age	(a1)	The client's age
job	(a2)	The client's job
marital	(a3)	The client's marital statue
education	(a4)	The client's highest completed level of education
default	(a5)	Whether the client's credit has defaulted or not
balance	(a6)	The client's average yearly balance
housing	(a7)	Whether the client has a housing loan or not
loan	(a8)	Whether the client has a personal loan or not
contact	(a9)	How the client was contacted
day	(a10)	The last day the client was contacted
month	(a11)	The last month the client was contacted
duration	(a12)	The duration of the last contact
campaign	(a13)	The number of times the client was contacted for this campaign
pdays	(a14)	The number of days that have passed since the client was contacted from a previous campaign
previous	(a15)	The number of times this client was contacted for previous campaigns
poutcome	(a16)	The outcome of previous marketing campaigns

Table 1: Description of attributes

An exploration of the dataset using a text editor and the Weka Explorer interface reveals the following...

- The number of clients (people who were contacted by the banking institution) in the dataset is 36,169.

- 31,981 clients did not subscribe to a term deposit, 4188 clients did subscribe to a term deposit.

- 88.4% of clients did not subscribe to a term deposit. This is the accuracy of the default classifier.

- There are 16 attributes (excluding the output attribute, 'termDeposit'), reducing the number of attributes may be beneficial to avoid overfitting..

- There are no missing values. However, some attributes have an 'unknown' value.

    - job (a2): Only 223 clients had an unknown job.

    - education (a4): Only 1480 clients had an unknown education.

    - contact (a9): 10417 clients were contacted by unknown means. The histogram (see Figure 1) shows that clients contacted by a cellular device (23416) were

more likely to make a deposit compared to clients contacted by telephone (2336), so it would have been useful to know how all clients were contacted

– poutcome (a16): The value of this is 'unknown' for 29621 clients. This abnormally large number may be due to the fact that 29616 clients were never contacted for previous campaigns.

- Viewing the histograms for each variable showed that clients who had defaulted would never subscribe to a term deposit

- Some attributes distribution of values were quite imbalanced.

  – balance (a6): Most clients had a low/negative balance.

  – month (a11): 25800 clients were contacted in May, June, July, and August.

  – pdays (a14): Very negatively skewed distribution, 29712 clients had a value of -1, meaning they were not previously contacted

  – previous (a15): Very negatively skewed distribution, 29616 clients were not contacted for previous marketing campaigns.

  – job (a2): There are 13 values but 7 of them include only 7267 clients

- Two-dimensional scatter plots do not show any strong class separation for any of the attributes. From this, we can infer that several attributes will be needed to determine whether a client will subscribe to a term deposit or not.


Figure 1: Histogram of the contact attribute

2.1.1 The Test Dataset

cworkPredict.arff will be used to evaluate the best-performing models on the training dataset which will be chosen using 10-fold cross validation. The number of clients in this dataset is 9042. 87.8% of the clients in this dataset did not subscribe to a term deposit, which is almost the same as the training set.

3 Data Preprocessing

The data provided for us was already split into a training dataset and an evaluation dataset. The test dataset and the training dataset both have a similar proportion of clients who did not subscribe to a term deposit.

The numeric attributes in the dataset are age (a1), balance (a6), day (a10), duration (a12), campaign (a13), pdays (a14), and previous (a15). During the initial exploration of the dataset, we discovered that pdays (a14) and previous (a15) were very negatively skewed. Most clients also had a low / negative balance (a6) (see Figure 2). These non-normal distributions may affect the naive Bayes classifier which assumes that numerical values have a normal distribution. To resolve this, we may need to discretise our numeric attributes. This will be tested when we train our models by using preprocessing filters.


Figure 2: Histogram of the balance attribute

4 Classification Models

We are now going to develop classification models using the training set and select the best performing models. This will be carried out in five steps: benchmark models, attribute selection, model development, combining models, model selection, and cost-based modelling.

4.1 Benchmark Models

The naive Bayes, k-nearest neighbour, J4.8, OneR, and logistic regression classifiers were applied with default parameters and without any attribute preprocessing on the training dataset. This gave us a basic benchmark for each of these classifiers. 10-fold cross-validation was used to increase the reliability of our estimates. For the k-nearest neighbour classifier, we enabled the cross-validation option and set the kNN parameter to 10. The resulting value of k was 9.

Model	Accuracy
Naive Bayes	88.1%
k-nearest neighbour (k = 9)	89.2%
J4.8	90.4%
OneR	88.5%
logistic regression	90.2%

Table 2: Basic benchmark results on training dataset

4.2 Attribute Selection

In the data exploration stage, we determined that finding a subset of important attributes may help prevent overfitting. Overfitting is when a model becomes too complex. It starts memorising data, including noise from the dataset. By reducing the number of attributes used by the classification model, overfitting will hopefully by avoided.

- By visualising the decision tree of the J4.8 classifier, we can see that the duration (a12) attribute is at the root node

- Attribute selection using the CFS subset evaluator picked five attributes: marital (a3), housing (a7), loan (a8), duration (a12), poutcome (a16).

- Attribute selection using the information gain ranking filter ranked the attributes in the following order: duration (a12), poutcome (a16), pdays (a14), month (a11), contact (a9), age (a1), previous (a15), housing (a7), job (a2), day (a10), balance (a6), campaign (a13), loan (a8), education (a4), marital (a3), and default (a5).

- Attribute selection using the gain ratio feature evaluator ranked the attributes in the following order: poutcome (a16), duration (a12), pdays (a14), previous (a15), contact (a9), housing (a7), month (a11), age (a1), loan (a8), balance (a6), job (a2), campaign (a13), default (a5), day (a10), marital (a3), and education (a4).

- Attribute selection using the symmetrical uncertainty ranking filter ranked the attributes in the following order: duration (a12), poutcome (a16), pdays (a14), previous (a15), contact (a9), month (a11), housing (a7), age (a1), balance (a6), job (a2), loan (a8), campaign (a13), day (a10), marital (a3), education (a4), default (a5).

- Attribute selection using the Chi-squared ranking filter ranked the attributes in the following order: duration (a12), poutcome (a16), pdays (a14), month (a11), age (a1), previous (a15), contact (a9), job (a2), housing (a7), day (a10), balance (a6), campaign (a13), education (a4), marital (a3), loan (a8), default (a5).

All of these attribute evaluators agree that the two most important attributes are duration (a12) and poutcome (a16). pdays (a14) is the third most important attribute according to all attribute evaluators apart from the CFS subset evaluator. Using these three attributes, the performance of the benchmark models is shown in Table 3. Reducing the number of attributes to 3 improves the performance of the naive Bayes and the k-nearest neighbour classifiers. The performance of OneR remains unchanged and the performance of J4.8 and logistic regression is slightly reduced.

Model	Accuracy
Naive Bayes	89.1%
k-nearest neighbour (k = 9)	89.7%
J4.8	90.1%
OneR	88.5%
logistic regression	90.0%

Table 3: Benchmark results on three-attribute dataset

The information gain ranking filter, the gain ratio feature evaluator, and the symmetrical uncertainty ranking filter all have the same attributes in their top eight but in a different order. The Chi-squared ranking filter also has seven of these eight attributes, with housing (a7), appearing in ninth place. Because of this, we have determined an intermediate selection of eight attributes to be: duration (a12), poutcome (a16), pdays (a14), month (a11), contact (a9), age (a1), previous (a15), and housing (a7). The performance of the classifiers with these eight attributes as seen in Table 4 show that increasing the number of attributes to eight slightly reduces the performance of the k-nearest neighbour classifier. The performance of the J4.8 and the logistic regression classifiers are slightly increased.

Model	Accuracy
Naive Bayes	89.1%
k-nearest neighbour (k = 9)	89.6%
J4.8	90.3%
OneR	88.5%
logistic regression	90.1%

Table 4: Benchmark results on eight-attribute dataset

4.3 Model Development

During this phase, we experiment with each of our classifiers to enable us to select the most optimal parameters for each model.

4.3.1 Naive Bayes

This model had the same performance on the three-attribute dataset and the eight-attribute dataset so we experimented with both datasets to determine the settings for the most optimal model. The naive Bayes classifier assumes that numeric attributes have a normal distribution. During our initial data exploration, we discovered that some attributes, namely pdays (a14) and previous (a15) were very negatively skewed. The balance (a6) attribute also had a negatively skewed distribution but this attribute is not in the datasets used to train this model. The easiest way to resolve this issue is to run the discretise filter. This will discretise numerical attributes into a specified number of nominal bins.

Attributes in both datasets were discretised into 10 bins of equal frequency. For the three-attribute dataset, the naive Bayes classifier achieved an accuracy of 89.3%, a

slight improvement when compared to the base benchmark. An accuracy of 88.2% was achieved by the naive Bayes classifier after discretising the eight-attribute dataset into 10 bins of equal frequency.

The three attribute dataset with 10 bins of equal frequency gives us the most accurate model for the naive Bayes classifier. Table 5 shows the accuracy of the naive Bayes classifier when different numbers of bins are used on the discretise filter.

Bins	Accuracy
5	89.2%
10	89.3%
15	89.4%
20	88.4%

Table 5: Naive Bayes accuracy on the three-attribute dataset when the numerical attributes are discretised into a specified number of bins

These results show that for the naive Bayes classifier, you get the best performing model by using the three-attribute dataset and discretising the numerical attributes into 15 bins of equal frequency.

4.3.2   k-nearest Neighbour

For this model, we continued to use 9 neighbours and the three-attribute dataset as the model performed best on this dataset. We experimented with the distance weighting parameters and the results can be found in Table 6

Weighting	Accuracy
No distance weighting	89.7%
1/distance	89.7%
1-distance	89.7%

Table 6: Results of varying the distance parameter

As you can see in Table 6, changing the distance weighting had no effect on the accuracy of this model. No distance weighting needs to be done.

4.3.3   J4.8

For the J4.8 classifier, the key parameters to experiment with are the ones that control the complexity of the model. These parameters are post-pruning, reduced error pruning, and minimum number of objects. J4.8 performed best with all attributes so all of the experiments were carried out on the original training dataset (see Table 7).

Parameter	Parameter Value	Accuracy
Confidence Factor	0.40	89.9%
Confidence Factor	0.35	90.1%
Confidence Factor	0.30	90.3%
Confidence Factor	0.25	90.4%
Confidence Factor	0.20	90.5%
Confidence Factor	0.15	90.5%
Confidence Factor	0.10	90.4%
Reduced error pruning (seed = 1)	True	90.1%
Minimum number of objects	10	90.6%
Minimum number of objects	20	90.3%
Minimum number of objects	30	90.4%

Table 7: Results of modifying J4.8 complexity parameters

The table above shows that the best performing model for this classifier is produced when the minimum number of objects parameter is set to 10.

4.3.4 OneR

All of the benchmark models for the OneR classifier had the same performance for all attribute sets. This is because the OneR classifier only uses one attribute for determining whether a client will subscribe to a term deposit or not. So, to allow the classifier to choose the best attribute for itself, we will use the full attribute set to further develop this model.

The OneR classifier discretises all numerical attributes into different buckets, similar to what the discretise attribute filter does. The classifier takes in a parameter, "minBucketSize", which specifies the minimum number of objects that must be in a bucket. In Table 8, we experimented with this parameter and determined that a minimum bucket size of 18 gives us the most accurate model.

Minimum Bucket Size	Accuracy
6	88.5%
8	88.6%
10	88.6%
12	88.7%
14	88.9%
16	89.2%
18	89.3%

Table 8: Results of varying the minimum bucket size parameter

4.3.5 Logistic Regression

The full training dataset gave us the best performance for this classifier so that is what we used. This classifier uses a ridge estimator and the ridge value is available for us to

alter as a parameter. A smaller ridge parameter results in a more flexible model. The results of experimenting with the ridge value are shown in Table 9

Ridge Value	Accuracy
1×10^{-8}	90.2%
1×10^{-4}	90.2%
1×1	90.2%
1×10	90.2%

Table 9: Results of varying the ridge value

As you can see in Table 9, changing the ridge parameter did not alter the accuracy of the resulting model in any meaningful way, so we decided to keep the ridge parameter at it's default setting (1×10^{-8})

4.3.6 Voting Committee

We now have five strongly performing models. All these models have an accuracy that is better than the default accuracy (88.4%). We combined all these models into a voting committee; all the models will vote on whether a client will subscribe to a term deposit or not and the majority wins. The models chosen are naive Bayes with numerical attributes being discretised into 15 bins of equal frequency, k-nearest neighbour with k set to 9, J4.8 with the minimum number of objects parameter set to 10, OneR with the minimum bucket size parameter set to 18, and logistic regression with default parameters.

Because of the difficulties of using a different training set for different models in a voting committee, we decided to train our models on the eight-attribute dataset. This dataset gives us the best performance for the naive Bayes and the OneR classifiers. The performances of the other classifiers are only very slightly reduced on the eight attribute dataset when compared to the optimal dataset for these classifiers.

After running the voting committee on the dataset with the best parameter settings for all five classifiers, t he overall performance with 10-fold cross-validation was **90.0%**. While this is pretty accurate, the voting committee is still beaten by both the logistic regression model and the J4.8 model.

4.3.7 Model Selection

By experimenting with different classifiers and their parameters, we have determined that the J4.8 classifier with the minimum number of objects parameter set to 10 produces the most accurate model. The J4.8 model has an estimated accuracy of 90.6% when using 10-fold cross-validation.

4.4 Cost-based Modelling

In this section, we ran all five models developed in the previous section using an unequal cross matrix. In the unequal cost scenario specified by the banking institution, the

cost of misclassifying a client who will subscribe to a term deposit is 10 times that of misclassifying a client who will not. The cost matrix for this has the following structure:

$$
\begin{matrix}
0.0 & 1.0 \\
10.0 & 0.0
\end{matrix}
$$

The default cost on the training set for a classifier that predicts that everyone will subscribe to a term deposit is equal to the number of people that will not subscribe to a term deposit; 31981.

4.4.1 Naive Bayes

We ran the naive Bayes classifier with the three-attribute dataset and numerical attributes discretised into 15 bins with the cost sensitive meta classifier. The naive Bayes model does give us estimates of class posterior probabilities so we will run it with the minimise expected cost parameter set to true. This gave us the following confusion matrix:

$$
\begin{matrix}
24791 & 7190 \\
961 & 3227
\end{matrix}
$$

The cost of this matrix is $7190 \times 1 + 961 \times 10 = 16800$.

4.4.2 k-nearest Neighbour

We ran this classifier with the cost sensitive meta classifier, using the three-attribute dataset, 9 neighbours, and no distance weighting. Because the k-nearest neighbour classifier does not give us estimates of class posterior probability, we set the minimise expected cost parameter to false. This tells the meta classifier to reweigh the training instances using the cost matrix. This resulted in the following confusion matrix:

$$
\begin{matrix}
23827 & 8154 \\
962 & 3226
\end{matrix}
$$

The cost of this matrix is $8154 \times 1 + 962 \times 10 = 17704$. This is more than the cost for naive Bayes.

4.4.3 J4.8

This classifier was ran with the cost sensitive meta classifier using the full training dataset and the minimum number of objects parameter set to 10. Like k-nearest neighbour, this classifier does not give us estimates of class posterior probabilities so we set the minimise expected cost parameter to false. This gave us the following confusion matrix:

$$\begin{matrix} 26270 & 5711 \\ 485 & 3703 \end{matrix}$$

The cost of this matrix is $5711 \times 1 + 485 \times 10 = 10561$. This is the best cost so far. margin.

4.4.4 OneR

For the OneR classifier, we used the full training set and set the minimum bucket size parameter to 18. This classifier also does not give us estimates of class posterior probability so we set the minimise expected cost parameter of the cost sensitive meta classifier to false. This gave us the following confusion matrix:

$$\begin{matrix} 20628 & 11353 \\ 990 & 3198 \end{matrix}$$

The cost of this matrix is $11353 \times 1 + 990 \times 10 = 21253$. This is the significantly worse than all other costs.

4.4.5 Logistic Regression

We used the full training set again for this classifier with default parameters. Since this classifier does give us posterior probabilities, we set the minimise expected cost parameter of the cost sensitive meta classifier to true. The confusion matrix produced is shown below:

$$\begin{matrix} 25375 & 6606 \\ 514 & 3674 \end{matrix}$$

The cost of this matrix is $6606 \times 1 + 514 \times 10 = 11746$. This is worse than the cost for the J4.8 model. other costs.

4.4.6 Model Selection

From the results reported here, we can clearly see that the J4.8 model is the best for the unequal cost problem. It significantly outperforms most of the other models, with only the logistic regression model coming close. Because of this, it is not worth developing a voting model, the J4.8 model is clearly the way to go

5 Conclusion

In this section, we will evaluate our best model, the J4.8 classifier for both the equal cost and the unequal cost scenario, on the test set. Then, I will talk about what I have learned by doing this coursework.

5.0.1 Evaluation

The accuracy of the J4.8 model with the minimum number of objects set to 10 on the test dataset with equal costs was 90.3%. This is pretty good, the accuracy of the default classifier on the test dataset 87.8%. The confusion matrix on the test set for this classifier was:

$$\begin{matrix} 7611 & 330 \\ 548 & 553 \end{matrix}$$

Running the cost sensitive J4.8 model on the test set with unequal costs gave us the following confusion matrix:

$$\begin{matrix} 6510 & 1431 \\ 121 & 980 \end{matrix}$$

The total cost of this confusion matrix is $1431 \times 1 + 121 \times 10 = 2641$. The cost of the default classifier would be equal to the number of clients who would not subscribe to a term deposit; 7941. So as we can see, the cost sensitive classifier performs very well on the test set.

5.0.2 What I have learned

By doing this coursework, I have reinforced the knowledge I gained during the lectures on the way the different classifiers work. I understand the preference to provide numerical attributes with a normal distribution for the naive Bayes classifier. I strengthened my understanding of how other classifiers like the J4.8 classifier, the logistic regression classifier, and the k-nearest neighbour classifier worked. Experimenting with the parameters of each of these classifiers gave me further insight into their inner workings, especially when I experimented with the number of bins parameter for the J4.8 classifier. I could clearly see how the J4.8 classifier discretises numerical attributes. I also grew an appreciation for simpler classifiers. The OneR classifier, even though it only contains one rule, was able to keep up with other, more complex classifiers.

I am much more comfortable with using Weka. This coursework enabled me to experiment with the attribute selector, showing me how to prune attributes which aren't very useful from a dataset. I also enjoyed creating a voting committee, I feel that using a voting committee with a lot of similarly performing models may be a very strong option in some scenarios.

Finally, I was able to identify holes in my knowledge by doing this coursework. This enabled me to go through previous labs, lectures, and online resources in order to reinforce my knowledge of machine learning and data mining.